

Q_2

November 2, 2023

1 PART A

```
[53]: import random

# Function to get the list of sum of the upward face value of n rolls for 1000
# repetitions
def roll_biased_die(k, n):
    results = []
    weights = [1 / (2 ** (k - 1))] + [1 / (2 ** (i - 1)) for i in range(2, k + 1)]
    for _ in range(1000):
        outcome = random.choices(range(1, k + 1), weights=weights, k=n)
        results.append(sum(outcome))

    return results
```

```
[54]: # To calculate probability with sum S for k faced dice rolled n times
def probability_sum(n, k, S):

    dp = [[0.0] * (S + 1) for _ in range(n + 1)]
    dp[0][0] = 1.0

    for i in range(1, n + 1):
        for j in range(1, S + 1):
            for face in range(1, k + 1):
                if j >= face:
                    if face == 1:
                        dp[i][j] += dp[i - 1][j - face] * (1 / (2 ** (k - 1)))
                    else:
                        dp[i][j] += dp[i - 1][j - face] * (1 / (2 ** (face - 1)))

    return dp
```

```
[55]: import numpy as np
import matplotlib.pyplot as plt
```

```

def fun(k, rolls):
    # Simulate die rolls and calculate the sum
    results = roll_biased_die(k, rolls)

    # Plot a frequency distribution histogram
    plt.hist(results)
    plt.xlabel("Sum of Upward Face Values")
    plt.ylabel("Frequency")
    plt.title(f"Frequency Distribution for {rolls} Rolls of a {k}-Faced Die")
    plt.show()

    # Calculate the theoretical expected sum
    dp = probability_sum(rolls, k, k*rolls)
    expected_sum_theoretical = 0
    for i in range(1, rolls*k+1):
        expected_sum_theoretical += i*dp[rolls][i]

    # Calculate the actual expected sum from the simulation
    expected_sum_actual = np.mean(results)

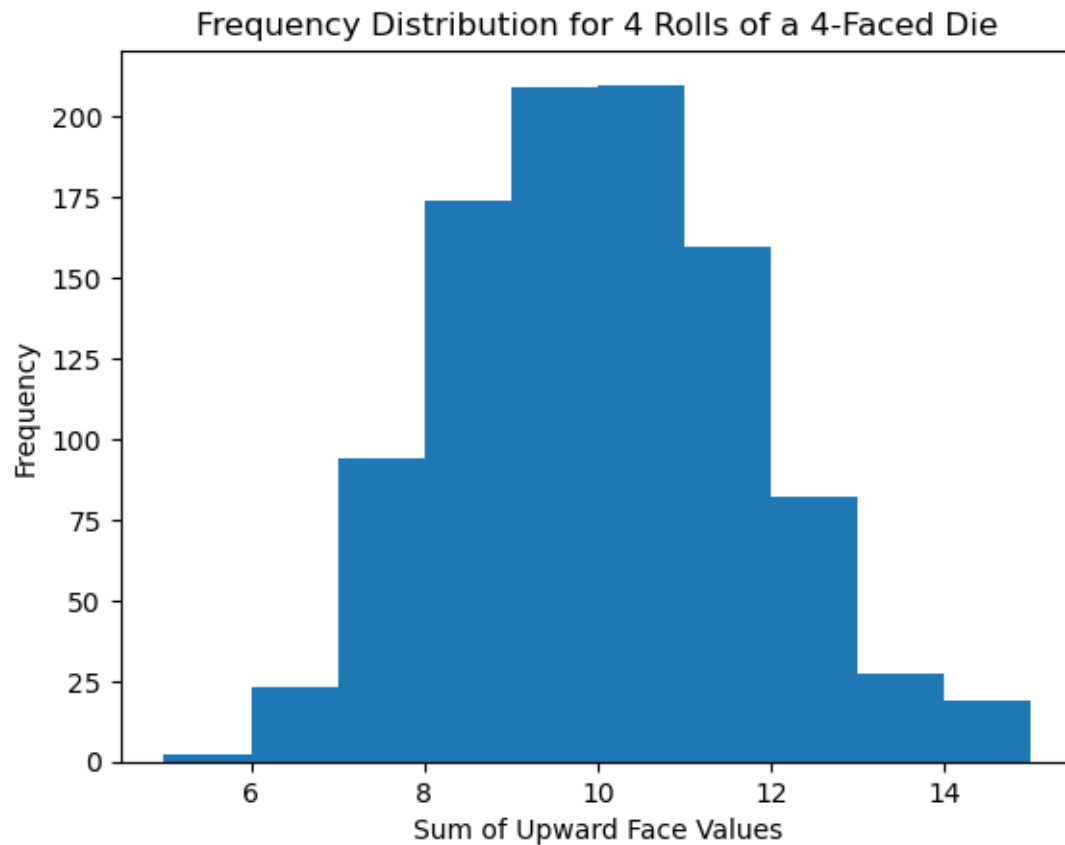
    print(f"Theoretical Expected Sum: {expected_sum_theoretical:.4f}")
    print(f"Actual Sum: {expected_sum_actual:.4f}")

    print("-----")

    # Print the five-number summary
    summary = np.percentile(results, [0, 25, 50, 75, 100])
    print(f"Minimum: {summary[0]}")
    print(f"1st Quartile: {summary[1]}")
    print(f"Median: {summary[2]}")
    print(f"3rd Quartile: {summary[3]}")
    print(f"Maximum: {summary[4]}")

```

```
[56]: fun(4,4)
```



Theoretical Expected Sum: 9.5000

Actual Sum: 9.5430

Minimum: 5.0


1st Quartile: 8.0

Median: 9.0

3rd Quartile: 11.0

Maximum: 15.0

[57]: `fun(4,8)`



Theoretical Expected Sum: 19.0000

Actual Sum: 19.0100

Minimum: 11.0


1st Quartile: 17.0

Median: 19.0

3rd Quartile: 21.0

Maximum: 27.0

[58]: `fun(16,4)`



Theoretical Expected Sum: 11.9979

Actual Sum: 12.0790

Minimum: 8.0

1st Quartile: 10.0

Median: 12.0

3rd Quartile: 14.0

Maximum: 26.0

2 PART B

2.0.1 Loading Dataset

```
[59]: from ucimlrepo import fetch_ucirepo
      from sklearn.model_selection import train_test_split

      spambase = fetch_ucirepo(id=94)
      X = spambase.data.features
      y = spambase.data.targets
      y = y.iloc[:, 0]
```


```
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3,
↪random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5,
↪random_state=42)
```

2.0.2 Plot Distribution of 5 columns

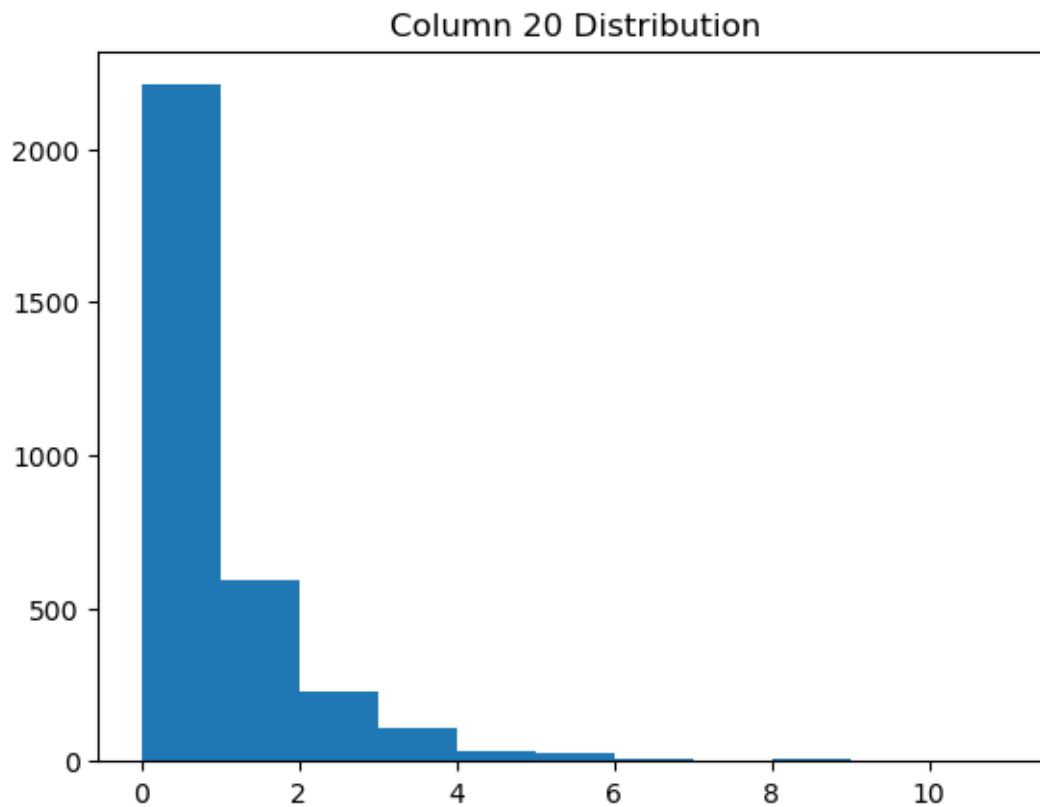
```
[60]: import matplotlib.pyplot as plt

mean_values = []
for col in range(0,54):
    a = X_train.to_numpy()[:, col]
    mean_values.append(np.mean(a))
indices = np.argsort(mean_values)[-5:] [::-1]
```


```
[61]: col = indices[0]
a = X_train.to_numpy()[:, col]
plt.hist(a,bins = range(0,12))
plt.title(f'Column {col} Distribution')
plt.show()
```




```
[62]: col = indices[1]
a = X_train.to_numpy()[:, col]
plt.hist(a, bins = range(0,12))
plt.title(f'Column {col} Distribution')
plt.show()
```




```
[63]: col = indices[2]
a = X_train.to_numpy()[:, col]
plt.hist(a, bins = range(0,8))
plt.title(f'Column {col} Distribution')
plt.show()
```



```
[64]: col = indices[3]
a = X_train.to_numpy()[:, col]
plt.hist(a, bins = range(0,12))
plt.title(f'Column {col} Distribution')
plt.show()
```

```
[65]: col = indices[4]
a = X_train.to_numpy()[:, col]
plt.hist(a, bins = range(0,8))
plt.title(f'Column {col} Distribution')
plt.show()
```



2.0.3 Priors of Classes

```
[66]: prior_spam = float(np.sum(y_train == 1) / len(y_train))
      prior_not_spam = float(np.sum(y_train == 0) / len(y_train))

      print(f'Prior Probability of Spam: {prior_spam}')
      print(f'Prior Probability of Not Spam: {prior_not_spam}')
```

Prior Probability of Spam: 0.3838509316770186
Prior Probability of Not Spam: 0.6161490683229813

2.0.4 Naive Bayes Classifier

```
[67]: class NaiveBayesClassifier():

      def calc_prior(self, features, target):

          self.prior = (features.groupby(target).apply(lambda x: len(x)) / self.
      ↪rows).to_numpy()

          return self.prior
```

```

def calc_statistics(self, features, target):

    self.mean = features.groupby(target).apply(np.mean, axis=0).to_numpy()
    self.var = features.groupby(target).apply(np.var, axis=0).to_numpy()

    return self.mean, self.var

def gaussian_density(self, class_idx, x):

    mean = self.mean[class_idx]
    var = self.var[class_idx]
    epsilon = 1e-10
    numerator = np.exp((-1/2)*((x-mean)**2) / (2 * (var + epsilon)))
    denominator = np.sqrt(2 * np.pi * (var + epsilon))
    prob = numerator / denominator
    return prob

def calc_posterior(self, x):
    posteriors = []

    # calculate posterior probability for each class
    for i in range(self.count):
        prior = np.log(self.prior[i]) ## use the log to make it more
        ↪ numerically stable
        conditional = np.sum(np.log(self.gaussian_density(i, x))) # use the
        ↪ log to make it more numerically stable
        posterior = prior + conditional
        posteriors.append(posterior)

    # return class with highest posterior probability
    return self.classes[np.argmax(posteriors)]

def total_parameters(self):
    # Calculate the total number of parameters needed to be stored for the
    ↪ model
    # Parameters to store include prior, mean, and variance
    total_parameters = self.count + self.count * self.feature_nums * 2
    return total_parameters

def predict_proba(self, features):
    # Initialize an empty array to store the class probabilities
    class_probabilities = np.zeros((features.shape[0], self.count))

    for i in range(self.count):
        # Calculate the log of the prior probability
        prior = np.log(self.prior[i])

```

```

        # Calculate the log of the conditional probability
        conditional = np.sum(np.log(self.gaussian_density(i, features)),
↪axis=1) # Sum along axis 1

        # Calculate the posterior probability (log scale)
        posterior = prior + conditional

        # Store the posterior probability for the current class
        class_probabilities[:, i] = posterior

        # Calculate class probabilities using the log-sum-exp trick for
↪numerical stability
        # This step converts log-probabilities to probabilities
        log_class_probabilities = class_probabilities - np.
↪max(class_probabilities, axis=1, keepdims=True)
        class_probabilities = np.exp(log_class_probabilities)

        # Normalize the probabilities to sum to 1 for each data point
        class_probabilities /= class_probabilities.sum(axis=1, keepdims=True)

        return class_probabilities

    def fit(self, features, target):
        self.classes = np.unique(target)
        self.count = len(self.classes)
        self.feature_nums = features.shape[1]
        self.rows = features.shape[0]

        self.calc_statistics(features, target)
        self.calc_prior(features, target)

    def predict(self, features):
        preds = [self.calc_posterior(f) for f in features.to_numpy()]
        return preds

```

2.0.5 Train the model

```

[68]: my_model = NaiveBayesClassifier()
      my_model.fit(X_train, y_train)

      print(f'Number of parameters to be stored: {my_model.total_parameters()}')

```

Number of parameters to be stored: 230

2.0.6 Evaluate the model

```
[69]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

def evaluate_model(classifier, X, y):
    y_pred = classifier.predict(X)
    accuracy = accuracy_score(y, y_pred)
    precision = precision_score(y, y_pred)
    recall = recall_score(y, y_pred)
    f1 = f1_score(y, y_pred)
    return accuracy, precision, recall, f1

accuracy, precision, recall, f1 = evaluate_model(my_model, X_test, y_test)
print("Naive Bayes Model Performance:")
print("-----")
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1-score: {f1}")
```

Naive Bayes Model Performance:

Accuracy: 0.8002894356005789

Precision: 0.6840796019900498

Recall: 0.9615384615384616

F1-score: 0.7994186046511628

C:\Users\user\AppData\Local\Temp\ipykernel_4856\3306703780.py:32:

RuntimeWarning: divide by zero encountered in log

conditional = np.sum(np.log(self.gaussian_density(i, x))) # use the log to
make it more numerically stable

2.0.7 Apply Log Transformation

```
[70]: X_train_log = X_train.apply(lambda x: np.log(x + 1))
X_val_log = X_val.apply(lambda x: np.log(x + 1))
X_test_log = X_test.apply(lambda x: np.log(x + 1))

# Evaluate the model with log-transformed data
accuracy_log, precision_log, recall_log, f1_log = evaluate_model(my_model,
    X_test_log, y_test)
print("Naive Bayes Model Performance (with log transformation):")
print("-----")
print(f"Accuracy: {accuracy_log}")
print(f"Precision: {precision_log}")
print(f"Recall: {recall_log}")
print(f"F1-score: {f1_log}")
```

Naive Bayes Model Performance (with log transformation):

```
-----  
Accuracy: 0.7264833574529667  
Precision: 0.6061269146608315  
Recall: 0.9685314685314685  
F1-score: 0.7456258411843876
```

```
C:\Users\user\AppData\Local\Temp\ipykernel_4856\3306703780.py:32:  
RuntimeWarning: divide by zero encountered in log  
    conditional = np.sum(np.log(self.gaussian_density(i, x))) # use the log to  
make it more numerically stable
```

2.0.8 Observation

After applying log transformation, the accuracy, precision, recall and f1-scores got reduced.

3 PART C

3.0.1 GaussianNB from sklearn

```
[71]: from sklearn.naive_bayes import GaussianNB  
  
gnb = GaussianNB(var_smoothing=1e-09)  
gnb.fit(X_train, y_train)  
  
accuracy, precision, recall, f1 = evaluate_model(gnb, X_test, y_test)  
print("Gaussian Naive Bayes Model Performance:")  
print("-----")  
print(f"Accuracy: {accuracy}")  
print(f"Precision: {precision}")  
print(f"Recall: {recall}")  
print(f"F1-score: {f1}")
```

Gaussian Naive Bayes Model Performance:

```
-----  
Accuracy: 0.8277858176555717  
Precision: 0.7191601049868767  
Recall: 0.958041958041958  
F1-score: 0.8215892053973014
```

```
[72]: # Evaluate the model with log-transformed data  
  
accuracy_log, precision_log, recall_log, f1_log = evaluate_model(gnb,  
    ↪ X_test_log, y_test)  
print("Gaussian Naive Bayes Model Performance (with log transformation):")  
print("-----")  
print(f"Accuracy: {accuracy_log}")  
print(f"Precision: {precision_log}")
```

```
print(f"Recall: {recall_log}")
print(f"F1-score: {f1_log}")
```

Gaussian Naive Bayes Model Performance (with log transformation):

```
-----
Accuracy: 0.788712011577424
Precision: 0.6699029126213593
Recall: 0.965034965034965
F1-score: 0.7908309455587392
```

3.0.2 ROC Curve

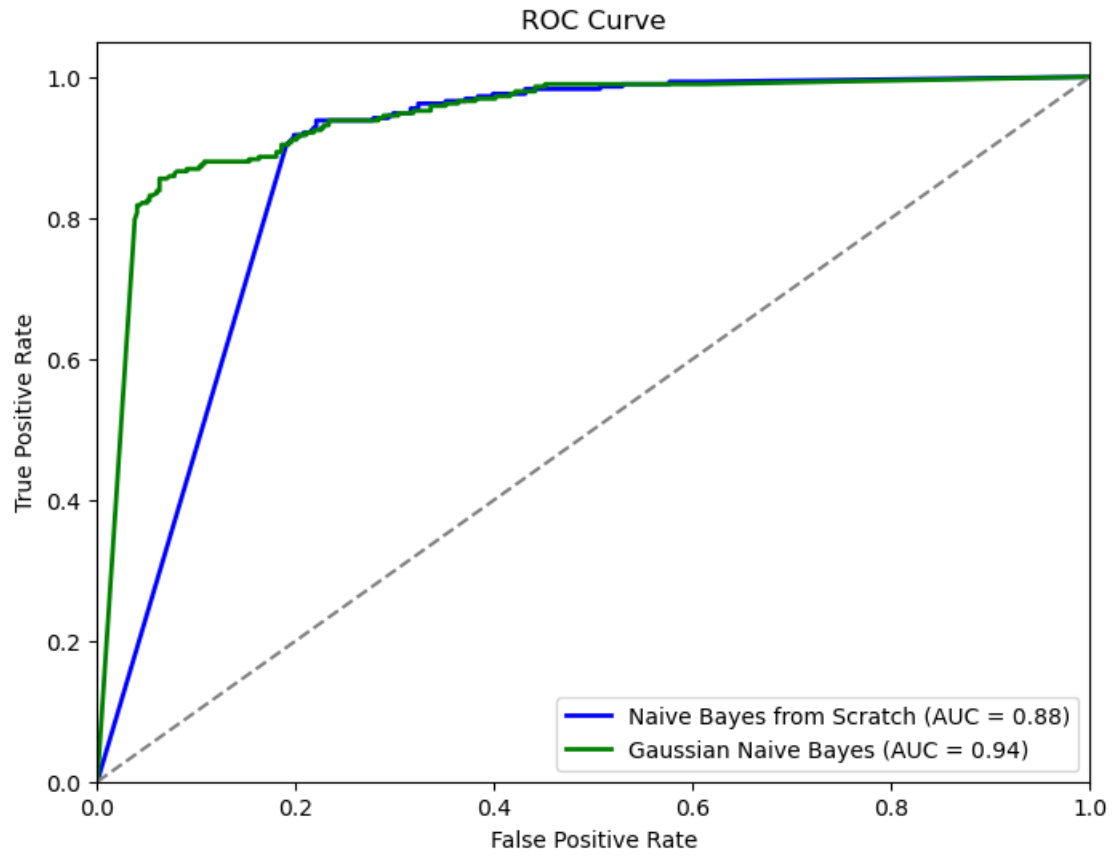
```
[73]: import matplotlib.pyplot as plt
      from sklearn.metrics import roc_curve, roc_auc_score

      # Calculate ROC curve and AUC for Naive Bayes from scratch
      y_scores_nb = my_model.predict_proba(X_val)[: , 1]
      fpr_nb, tpr_nb, _ = roc_curve(y_val, y_scores_nb)
      roc_auc_nb = roc_auc_score(y_val, y_scores_nb)

      # Calculate ROC curve and AUC for Gaussian Naive Bayes from scikit-learn
      y_scores_gnb = gnb.predict_proba(X_val)[: , 1]
      fpr_gnb, tpr_gnb, _ = roc_curve(y_val, y_scores_gnb)
      roc_auc_gnb = roc_auc_score(y_val, y_scores_gnb)

      # Plot the ROC curves
      plt.figure(figsize=(8, 6))
      plt.plot(fpr_nb, tpr_nb, color='blue', lw=2, label=f'Naive Bayes from Scratch_
      ↳(AUC = {roc_auc_nb:.2f})')
      plt.plot(fpr_gnb, tpr_gnb, color='green', lw=2, label=f'Gaussian Naive Bayes_
      ↳(AUC = {roc_auc_gnb:.2f})')
      plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
      plt.xlim([0.0, 1.0])
      plt.ylim([0.0, 1.05])
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve')
      plt.legend(loc='lower right')
      plt.show()
```

```
C:\Users\user\AppData\Roaming\Python\Python311\site-
packages\pandas\core\internals\blocks.py:351: RuntimeWarning: divide by zero
encountered in log
      result = func(self.values, **kwargs)
```



3.0.3 Observations

Higher AUC for Gaussian Naive Bayes than the Naive Bayes Implemented. So GNB is better than Naive Bayes implemented for email spam classification.